

Applying Pose Estimation to Predict the Outcome of Basketball Shots

Laith Altarabishi

University of Texas at Austin

laithaustin@utexas.edu

Sidharth Babu

University of Texas at Austin

sidharth.n.babu@utexas.edu

Afnan Mir

University of Texas at Austin

afnanmir@utexas.edu

Zayam Tariq

University of Texas at Austin

zayamtariq@utexas.edu

Abstract

In this report, we present an end-to-end pipeline that includes a pose estimator as well as a classification model that can effectively predict the outcome of a free throw. Our training data was self-generated by recording hundreds of clips of our test subject shooting free throws of various forms and recording the outcome. We explore different ways to generate feature vectors for inputs to our model as well as multiple classification models to produce the best performing pipeline.

1. Introduction

In basketball, the ability of a player to effectively shoot the basketball typically comes down to the player's shooting form. While the form of the best shooters tend to look different, they all typically use the same fundamentals. In our project, we will attempt to capture these fundamental aspects of a player's shooting form and attempt to predict the outcome of a shot using these features. Research has been done on extracting features from a player's movement to classify the action a player is performing (shooting, dribbling, etc.) [7], but we would like to focus our energy on feature extraction from the shooting motion using pose estimation [1], object detection [11], and possibly other methods to extract feature descriptors of a shot and attempt to identify it as a make or a miss.

This problem is a particularly nontrivial application of pose estimation for two main reasons. The first being that there are multiple stages to a basketball shot that need to be taken into account. From dipping the ball to waist level, to the motion of bringing the ball to eye level, to releasing the ball, each plays a significant role in the outcome of a shot, so each stage needs to be taken into account. The second reason is the variability of the average shot. It can be argued that no two shots will ever be identical due to the

imperfect nature of humans. Therefore, it is necessary for our feature representation of a shot to be invariant towards miniscule changes in shot form and focus more on fundamental differences.

2. Related Works

2.1. Pose Estimation

Pose Estimation is a critical topic in computer vision that will intersect with our goal of trying to accurately capture the motion and actions of a person taking a shot in basketball. Pose estimation, in the context of 2D videos of humans, is the problem of localizing anatomical keypoints or joints in a frame by frame video or image [1]. To fulfill our goal of predicting the outcome of a basketball shot, it will be critical to assess the form of a player who's taking a shot - where form can be decomposed into various classifications of joints in space. Pose estimation methods can be categorized into bottom-up or top-down methodologies. Bottom-up methodologies start by estimating keypoints and body joints first, and then these points are clustered to form poses. In contrast, top-down methodologies of pose estimation first run a person detector before decomposing each person into their respective body joints within detected bounding boxes [14]. Computational complexity is a major consideration for landmark pose estimation algorithms, and modern SOTA pose estimation algorithms deploy deep learning and CNNs to improve computational overhead and speed [1]. We list some examples of prevalent and SOTA pose estimation models that have been employed and researched below.

OpenPose: The first multi-person realtime 2D pose estimation system that uses a bottom-up approach that implements nonparametric representation to associate human keypoints and body parts with an individual in an image [1].

DeepPose: SOTA pose estimation method that uses DNNs to classify human body joints through the usage of cascading DNN regressors that produce high precision pose estimates [12].

AlphaPose: Multi-person SOTA realtime pose estimation system that outperforms OpenPose in AP score and has a high mAP score [5].

DeepCut: Proposes an approach to solving issues in both pose estimation and detection by using a partitioning and labeling formulation of a set of CNN part detectors [10].

Pose estimation attempts to detect the location of 17 keypoints on a human body, as seen in Figure 1. These keypoints include key joints such as the knee, elbow, shoulder, and wrist, which can be vital in determining the form of a free throw.




Figure 1. Pose estimation keypoints

2.2. Basketball and Pose Estimation

The application of pose estimation in basketball is not a new concept. Collecting and analyzing basketball player's posture data is an important facet of the scientific basketball community in order to help maximize training outputs. For example, pose estimation was used in combination with classification models to predict the action a player is performing in a video [7]. This is fairly similar to our project because it attempts to create feature vectors in video frames using pose estimation to encode data about the motion of a player and use these features to make a prediction. However, our project solely focuses on the motion of a player's shooting form and predicting the outcome of the shot.

2.3. Object Detection

Our project will hope to capture and detect objects in high frame-rate video, with minimal computational over-

head so that we are able to best assess the keypoints/human joints of our basketball shooters. In recent years, object detection algorithms have evolved greatly and most SOTA models today utilize deep learning to provide more robust results [13]. There are many pre-existing SOTA object detection methods that have been utilized in the context of pose estimation methods and regression issues, and the following works are examples of some of them.

YOLO: Single stage object detection algorithm that frames detection as a regression problem to spatially separated bounding boxes and associated class probabilities [11].

Mask R-CNN: Two-stage object detection algorithm that detects objects in an image while creating high-quality segmentation masks for each instance [6].

Feature Pyramid Networks: Two-stage object detection algorithm that uses the multi-scale, pyramidal hierarchy of deep convolutional networks to create feature pyramids [8].

3. Methodology

3.1. General Overview

The high-level overview of our end-to-end pipeline is shown in Figure 2. We start off by capturing clips of a test subject shooting a free throw. This will be our original data. We first want to downsample the video as we assume that many of the frames will be redundant and take up unnecessary computation time. After this, we perform pose estimation on all the frames of each video to get the keypoints of our test subject for each video clip. Using the pose estimation data, we perform feature extraction to generate feature vectors for each video clip according to some convention. This will be elaborated on in a later section.

These feature vectors will be the training and testing data that we use for our classification model. We will want to split this data into training and testing sets. We will train our classification model on the training set and generate predictions using our testing set.

3.2. Data Collection

To fit the scope of our project, we wanted the shot data be as controlled as possible to prevent unwanted variations in our data. To do this, we collected our own data, where we had our test subject always shoot from the free throw line and had a camera in a fixed position in front of the test subject to capture their form. The set up can be seen in Figure 3. We generated approximately 300 clips of our test subject shooting free throws. Additionally, for this project, we

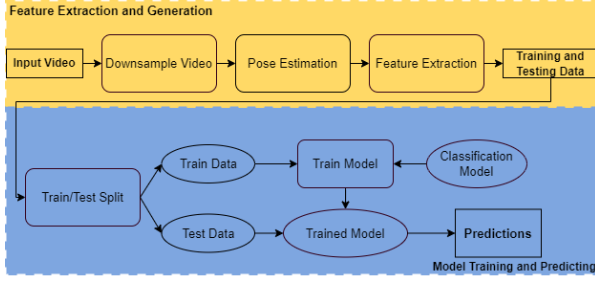


Figure 2. Overview of Pipeline



Figure 3. Free Throw

wanted to detect major changes in shooting forms that could affect the outcome of a shot. Therefore, our test subject shot some of their shots as he normally would, with the intent of making it, and some of their shots with purposefully bad form, with the intent of missing. We had a relatively even split of these two types of shots.

3.3. Pose Estimation

For our pose estimation, we decided between using OpenPose and MoveNet pose estimators. These are both bottom up approaches to pose estimation. We decided on using bottom up pose estimators as this is the current state of the art, and it is computationally less expensive [2]. Our first choice was OpenPose, which is considered the SOTA pose estimator. However, we found that OpenPose had some issues with our data. We found that having the basketball greatly hindered the performance of OpenPose detecting our test subject, as seen in Figure 4. Additionally, we found OpenPose to be relatively slow on the edge compared to MoveNet. In contrast, MoveNet was not hindered as much by the basketball and was much faster on the edge. Therefore, we decided to move forward with MoveNet as our pose estimator.

3.4. Feature Vector Generation

Before we can train a classification model to predict the outcome of a shot in a video clip, we need to define how we



Figure 4. OpenPose Pose Detection




Figure 5. MoveNet Pose Detection

are going to generate the feature vectors for each video clip.

As stated previously, we first want to downsample the video to reduce the number of frames we are processing, as many of the frames will be redundant. Our initial approach was to downsample the video by taking every n th frame, where n is a hyperparameter. However, we soon realized that our clips were not all the same length, which would mean we would receive an inconsistent number of frames per clip, which could potentially lead to dimensionality issues when we created our feature vectors, depending on our approach. To solve this, we decided to take an absolute number of frames per clip, which we decided to be 60.

After downsampling, we perform pose estimation on each frame of the video clip. Using the pose estimation data, we defined two possible approaches to feature generation:

1. Concatenate the pose estimation data from each frame of the clip into a single feature vector.
2. Use the pose estimation data from each frame of the clip to generate a feature vector for each frame, and perform majority voting the feature vectors to generate the final prediction.

Our pose estimator generates a set of 17 two dimensional coordinates, $S = \{(x_1, y_1), (x_2, y_2), \dots, (x_{17}, y_{17})\}$, where each coordinate (x_i, y_i) is the position on the image of keypoint i . For both approaches on each frame i ,




Figure 6. First Approach to Feature Vector Generation

we concatenate all the coordinates to create a feature vector $v_{ij} \in \mathbb{R}^{34}$, where v_{ij} is a feature vector for frame j of clip i . For our first approach, we take all the feature vectors for a single clip and concatenate them together, as seen in Figure 6. This takes 60 feature vectors of 34 dimensions and creates a feature vector $v_i \in \mathbb{R}^{2040}$, where v_i is the feature vector for clip i . In contrast, for our second approach, we treat each feature vector as a separate data point, and we perform a form of majority voting at the classification stage to generate our final prediction. Both of these approaches were tested, and performance scores were generated.

3.5. Classification Model

After we have generated our feature vectors, we used these vectors to train a classification model. We tested various classification models such as Gradient Boosted Trees (Catboost), Support Vector Machines, Multi-Layer Perceptrons, and K Nearest Neighbors [3] [4] [9]. We did preliminary testing on the classification models by measuring the performance of each model on our first feature vector approach as a means of choosing the best model for us to do a deeper dive on. To measure the performance of each model, we used accuracy, as generating the AUC score was not feasible for some of the models. Table 1 shows us the results of each model we tested:

Table 1. Accuracy Scores for Classification Models

Model	Accuracy
Catboost Gradient Boosting Tree	0.636
Support Vector Machine	0.580
Multi-Layer Perceptron	0.556
K Nearest Neighbors	0.602

From the results, we see that the Catboost model got the highest accuracy score. As a result, we chose this model to perform the rest of our experiments on.

References

- [1] Zhe Cao, Gines Hidalgo, Tomas Simon, Shih-En Wei, and Yaser Sheikh. Openpose: Realtime multi-person 2d pose estimation using part affinity fields, 2018. 1
- [2] Datagen. Pose estimation: Concepts, techniques, and how to get started, 2020. 3
- [3] Anna Veronika Dorogush, Andrey Gulin, Gleb Gusev, Nikita Kazeev, Liudmila Ostroumova Prokhorenkova, and Aleksandr Vorobev. Fighting biases with dynamic boosting. *CoRR*, abs/1706.09516, 2017. 4
- [4] Theodoros Evgeniou and Massimiliano Pontil. Support vector machines: Theory and applications. volume 2049, pages 249–257, 09 2001. 4
- [5] Haoshu Fang, Shuqin Xie, and Cewu Lu. RMPE: regional multi-person pose estimation. *CoRR*, abs/1612.00137, 2016. 2
- [6] Kaiming He, Georgia Gkioxari, Piotr Dollár, and Ross Girshick. Mask r-cnn, 2017. 2
- [7] Rong Ji. Research on basketball shooting action based on image feature extraction and machine learning. *IEEE Access*, 8:138743–138751, 2020. 1, 2
- [8] Tsung-Yi Lin, Piotr Dollár, Ross Girshick, Kaiming He, Bharath Hariharan, and Serge Belongie. Feature pyramid networks for object detection, 2016. 2
- [9] Fionn Murtagh. Multilayer perceptrons for classification and regression. *Neurocomputing*, 2(5):183–197, 1991. 4
- [10] Leonid Pishchulin, Eldar Insafutdinov, Siyu Tang, Bjoern Andres, Mykhaylo Andriluka, Peter Gehler, and Bernt Schiele. Deepcut: Joint subset partition and labeling for multi person pose estimation, 2015. 2
- [11] Joseph Redmon, Santosh Divvala, Ross Girshick, and Ali Farhadi. You only look once: Unified, real-time object detection. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 779–788, 2016. 1, 2
- [12] Alexander Toshev and Christian Szegedy. DeepPose: Human pose estimation via deep neural networks. In *2014 IEEE Conference on Computer Vision and Pattern Recognition*. IEEE, jun 2014. 2
- [13] Viso.ai. Object detection in 2022: The definitive guide, 2020. 2
- [14] Viso.ai. Human pose estimation with deep learning the ultimate overview, 2022. 1