

Applying Pose Estimation to Predict the Outcome of Basketball Shots

Laith Altarabishi

University of Texas at Austin

laithaustin@utexas.edu

Sidharth Babu

University of Texas at Austin

sidharth.n.babu@utexas.edu

Afnan Mir

University of Texas at Austin

afnanmir@utexas.edu

Zayam Tariq

University of Texas at Austin

zayamtariq@utexas.edu

Abstract

In this report, we present an end-to-end pipeline that includes a pose estimator as well as a classification model that can effectively predict the outcome of a free throw. Our training data was self-generated by recording hundreds of clips of our test subject shooting free throws of various forms and recording the outcome. We explore different ways to generate feature vectors for inputs to our model as well as multiple classification models to produce the best performing pipeline.

1. Introduction

In basketball, the ability of a player to effectively shoot the basketball typically comes down to the player's shooting form. While the form of the best shooters tend to look different, they all typically use the same fundamentals. In our project, we will attempt to capture these fundamental aspects of a player's shooting form and attempt to predict the outcome of a shot using these features. Research has been done on extracting features from a player's movement to classify the action a player is performing (shooting, dribbling, etc.) [7], but we would like to focus our energy on feature extraction from the shooting motion using pose estimation [1], object detection [11], and possibly other methods to extract feature descriptors of a shot and attempt to identify it as a make or a miss.

This problem is a particularly nontrivial application of pose estimation for two main reasons. The first being that there are multiple stages to a basketball shot that need to be taken into account. From dipping the ball to waist level, to the motion of bringing the ball to eye level, to releasing the ball, each plays a significant role in the outcome of a shot, so each stage needs to be taken into account. The second reason is the variability of the average shot. It can be argued that no two shots will ever be identical due to the

imperfect nature of humans. Therefore, it is necessary for our feature representation of a shot to be invariant towards miniscule changes in shot form and focus more on fundamental differences.

2. Related Works

2.1. Pose Estimation

Pose Estimation is a critical topic in computer vision that will intersect with our goal of trying to accurately capture the motion and actions of a person taking a shot in basketball. Pose estimation, in the context of 2D videos of humans, is the problem of localizing anatomical keypoints or joints in a frame by frame video or image [1]. To fulfill our goal of predicting the outcome of a basketball shot, it will be critical to assess the form of a player who's taking a shot - where form can be decomposed into various classifications of joints in space. Pose estimation methods can be categorized into bottom-up or top-down methodologies. Bottom-up methodologies start by estimating keypoints and body joints first, and then these points are clustered to form poses. In contrast, top-down methodologies of pose estimation first run a person detector before decomposing each person into their respective body joints within detected bounding boxes [14]. Computational complexity is a major consideration for landmark pose estimation algorithms, and modern SOTA pose estimation algorithms deploy deep learning and CNNs to improve computational overhead and speed [1]. We list some examples of prevalent and SOTA pose estimation models that have been employed and researched below.

OpenPose: The first multi-person realtime 2D pose estimation system that uses a bottom-up approach that implements nonparametric representation to associate human keypoints and body parts with an individual in an image [1].

DeepPose: SOTA pose estimation method that uses DNNs to classify human body joints through the usage of cascading DNN regressors that produce high precision pose estimates [12].

AlphaPose: Multi-person SOTA realtime pose estimation system that outperforms OpenPose in AP score and has a high mAP score [5].

DeepCut: Proposes an approach to solving issues in both pose estimation and detection by using a partitioning and labeling formulation of a set of CNN part detectors [10].

Pose estimation attempts to detect the location of 17 keypoints on a human body, as seen in Figure 1. These keypoints include key joints such as the knee, elbow, shoulder, and wrist, which can be vital in determining the form of a free throw.




Figure 1. Pose estimation keypoints

2.2. Basketball and Pose Estimation

The application of pose estimation in basketball is not a new concept. Collecting and analyzing basketball player's posture data is an important facet of the scientific basketball community in order to help maximize training outputs. For example, pose estimation was used in combination with classification models to predict the action a player is performing in a video [7]. This is fairly similar to our project because it attempts to create feature vectors in video frames using pose estimation to encode data about the motion of a player and use these features to make a prediction. However, our project solely focuses on the motion of a player's shooting form and predicting the outcome of the shot.

2.3. Object Detection

Our project will hope to capture and detect objects in high frame-rate video, with minimal computational over-

head so that we are able to best assess the keypoints/human joints of our basketball shooters. In recent years, object detection algorithms have evolved greatly and most SOTA models today utilize deep learning to provide more robust results [13]. There are many pre-existing SOTA object detection methods that have been utilized in the context of pose estimation methods and regression issues, and the following works are examples of some of them.

YOLO: Single stage object detection algorithm that frames detection as a regression problem to spatially separated bounding boxes and associated class probabilities [11].

Mask R-CNN: Two-stage object detection algorithm that detects objects in an image while creating high-quality segmentation masks for each instance [6].

Feature Pyramid Networks: Two-stage object detection algorithm that uses the multi-scale, pyramidal hierarchy of deep convolutional networks to create feature pyramids [8].

3. Methodology

3.1. General Overview

The high-level overview of our end-to-end pipeline is shown in Figure 2. We start off by capturing clips of a test subject shooting a free throw. This will be our original data. We first want to downsample the video as we assume that many of the frames will be redundant and take up unnecessary computation time. After this, we perform pose estimation on all the frames of each video to get the keypoints of our test subject for each video clip. Using the pose estimation data, we perform feature extraction to generate feature vectors for each video clip according to some convention. This will be elaborated on in a later section.

These feature vectors will be the training and testing data that we use for our classification model. We will want to split this data into training and testing sets. We will train our classification model on the training set and generate predictions using our testing set.

3.2. Data Collection

To fit the scope of our project, we wanted the shot data be as controlled as possible to prevent unwanted variations in our data. To do this, we collected our own data, where we had our test subject always shoot from the free throw line and had a camera in a fixed position in front of the test subject to capture their form. The set up can be seen in Figure 3. We generated approximately 300 clips of our test subject shooting free throws. Additionally, for this project, we

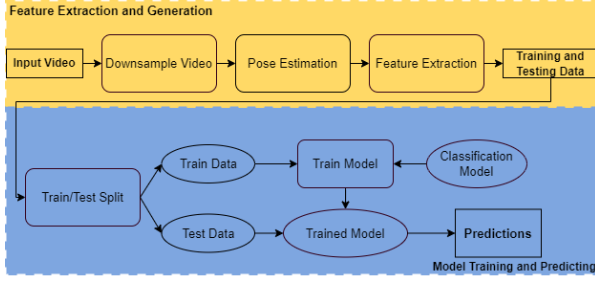


Figure 2. Overview of Pipeline



Figure 3. Free Throw

wanted to detect major changes in shooting forms that could affect the outcome of a shot. Therefore, our test subject shot some of their shots as he normally would, with the intent of making it, and some of their shots with purposefully bad form, with the intent of missing. We had a relatively even split of these two types of shots.

3.3. Pose Estimation

For our pose estimation, we decided between using OpenPose and MoveNet pose estimators. These are both bottom up approaches to pose estimation. We decided on using bottom up pose estimators as this is the current state of the art, and it is computationally less expensive [2]. Our first choice was OpenPose, which is considered the SOTA pose estimator. However, we found that OpenPose had some issues with our data. We found that having the basketball greatly hindered the performance of OpenPose detecting our test subject, as seen in Figure 4. Additionally, we found OpenPose to be relatively slow on the edge compared to MoveNet. In contrast, MoveNet was not hindered as much by the basketball and was much faster on the edge. Therefore, we decided to move forward with MoveNet as our pose estimator.

3.4. Feature Vector Generation

Before we can train a classification model to predict the outcome of a shot in a video clip, we need to define how we



Figure 4. OpenPose Pose Detection




Figure 5. MoveNet Pose Detection

are going to generate the feature vectors for each video clip.

As stated previously, we first want to downsample the video to reduce the number of frames we are processing, as many of the frames will be redundant. Our initial approach was to downsample the video by taking every n th frame, where n is a hyperparameter. However, we soon realized that our clips were not all the same length, which would mean we would receive an inconsistent number of frames per clip, which could potentially lead to dimensionality issues when we created our feature vectors, depending on our approach. To solve this, we decided to take an absolute number of frames per clip, which we decided to be 60.

After downsampling, we perform pose estimation on each frame of the video clip. Using the pose estimation data, we defined two possible approaches to feature generation:

1. Concatenate the pose estimation data from each frame of the clip into a single feature vector.
2. Use the pose estimation data from each frame of the clip to generate a feature vector for each frame, and perform majority voting the feature vectors to generate the final prediction.

Our pose estimator generates a set of 17 two dimensional coordinates, $S = \{(x_1, y_1), (x_2, y_2), \dots, (x_{17}, y_{17})\}$, where each coordinate (x_i, y_i) is the position on the image of keypoint i . For both approaches on each frame i ,




Figure 6. First Approach to Feature Vector Generation

we concatenate all the coordinates to create a feature vector $v_{ij} \in \mathbb{R}^{34}$, where v_{ij} is a feature vector for frame j of clip i . For our first approach, we take all the feature vectors for a single clip and concatenate them together, as seen in Figure 6. This takes 60 feature vectors of 34 dimensions and creates a feature vector $v_i \in \mathbb{R}^{2040}$, where v_i is the feature vector for clip i . In contrast, for our second approach, we treat each feature vector as a separate data point, and we perform a form of majority voting at the classification stage to generate our final prediction. Both of these approaches were tested, and performance scores were generated.

3.5. Classification Model

After we have generated our feature vectors, we used these vectors to train a classification model. We tested various classification models such as Gradient Boosted Trees (Catboost), Support Vector Machines, Multi-Layer Perceptrons, and K Nearest Neighbors [3] [4] [9]. We did preliminary testing on the classification models by measuring the performance of each model on our first feature vector approach as a means of choosing the best model for us to do a deeper dive on. To measure the performance of each model, we used accuracy, as generating the AUC score was not feasible for some of the models. Table 1 shows us the results of each model we tested:

Table 1. Accuracy Scores for Classification Models

Model	Accuracy
Catboost Gradient Boosting Tree	0.636
Support Vector Machine	0.580
Multi-Layer Perceptron	0.556
K Nearest Neighbors	0.602

From the results, we see that the Catboost model got the highest accuracy score. As a result, we chose this model to perform the rest of our experiments on.

4. Results

In this section, we present and discuss in detail the results of our experiments. In all of these experiments, we used MoveNet as our pose estimator and Catboost Gradient Boosting Tree as our classification model, as explained in the previous section.

4.1. First Approach to Feature Vector Generation

In our first experiment, we did a more extensive test on our first approach to feature vector generation. In our preliminary test, we did not use our full set of data, as we generated more data after this preliminary experiment. Additionally, we would like to evaluate the model on ROC-AUC score. Lastly, as we repeatedly retrained and reevaluated the model, we were getting performance scores that varied to some extent. Because of this, we decided to create a sample distribution of ROC-AUC scores we achieved for each iteration of our model. To do this, we ran our model 100 times, each with a randomized train-test split. We then recorded each ROC-AUC score and plotted the histogram of the scores. Figure 7 shows the distribution of our ROC-AUC scores for our first approach to feature generation.




Figure 7. Distribution of ROC-AUC Scores for First Approach to Feature Vector Generation

We produced a mean ROC-AUC score of 0.638, with a standard deviation of 0.043. We also produced a histogram distribution of our accuracy scores, and this distribution can be seen in Figure 8. We produced a mean accuracy score of 0.570, with a standard deviation of 0.044.

From these results, we see a performance that is better than random guessing, but not by much. Additionally, there is some nontrivial amount of deviation in the performance.

4.2. Second Approach to Feature Vector Generation

In our second experiment, we performed extensive performance tests on our second approach to feature generation. Firstly, we wanted to see how effective our classification model would be classifying each frame of every video




Figure 8. Distribution of Accuracy Scores for First Approach to Feature Vector Generation




Figure 9. Distribution of Accuracy Scores for Frame Classification

clip independently as a make or miss. This meant we had to label each feature vector for each frame as a make or miss, depending on what the outcome of the clip it originated from. We then trained our classification model on this data and evaluated its performance. Like we did before, we ran our model 100 times, each with a randomized train-test split and generated a distribution of accuracy scores. We can see the results in Figure 9. From this distribution, we can see that our classification model does fairly well in classifying each frame as part of a make clip or a miss clip.

With this in mind, we decided that we could improve performance of predicting the outcome of each clip by training the model on singular frames and producing our final prediction by performing some form of a majority vote on each frame for a given clip. This means, to predict the outcome of a certain clip, we would input the 60 associated feature vectors into the model for each frame. We predict the outcome for each of these frames, and if at least half of the frames predict a make, our final output will be a make. If at least half of the frames predict a miss, our final output will be a miss.

We ran this pipeline 100 times to generate a distribution of ROC-AUC scores and accuracy scores. These can be seen in Figure 10 and Figure 11 respectively.




Figure 10. Distribution of ROC-AUC Scores for Equal Weighted Majority Voting




Figure 11. Distribution of Accuracy Scores for Equal Weighted Majority Voting

We achieved a mean ROC-AUC score of 0.794 with a standard deviation of 0.063 and a mean accuracy score of 0.726 with a standard deviation of 0.068. Though the standard deviation is slightly higher, the overall performance of this approach seems to be much better than our first approach, with both ROC-AUC and accuracy scores increasing significantly.

4.3. Tuning the Majority Voting Threshold

After evaluating majority voting with equal weights, we wanted to see if we could improve the performance of this by tuning the threshold of our majority vote. We tried three different thresholds: 0.25, 0.5, and 0.75. For each threshold we, again, ran 100 iterations and generated a distribution of ROC-AUC scores and accuracy scores. The statistics of the

Threshold	Mean of ROC-AUC	Standard Deviation of ROC-AUC
0.25	0.795	0.067
0.50	0.794	0.063
0.75	0.793	0.058

Table 2. Statistics of ROC-AUC Majority Voting Threshold Experiments

Threshold	Mean of Accuracy	Standard Deviation of Accuracy
0.25	0.679	0.072
0.50	0.726	0.068
0.75	0.688	0.062

Table 3. Statistics of Accuracy Majority Voting Threshold Experiments

distributions for each of the thresholds can be seen in Table 2 and Table 3.

We see that the mean ROC-AUC score is highest for a threshold of 0.25, but only by a very small margin. In general, the ROC-AUC scores are quite similar, but it is clear that when measuring performance by accuracy, the threshold of 0.5 tends to perform the best. The combination of the ROC-AUC scores and accuracy scores lead us to believe that the threshold of 0.5 is the best choice.

4.4. Modifying Majority Voting Criteria

Our final experiment was to see if we could improve performance of the majority voting by using the median of the soft probability scores instead of the mean. This was done in an attempt to reduce the effect of outliers on the vote. Again, we ran this 100 times and generated a performance score distribution. We attempted to use three different thresholding values for this as well. The statistics of the ROC-AUC scores and accuracy scores

Threshold	Mean of ROC-AUC	Standard Deviation of ROC-AUC
0.25	0.783	0.057
0.50	0.798	0.053
0.75	0.793	0.059

Table 4. Statistics of ROC-AUC Median Voting Threshold Experiments

Threshold	Mean of Accuracy	Standard Deviation of Accuracy
0.25	0.688	0.057
0.50	0.733	0.054
0.75	0.663	0.078

Table 5. Statistics of Accuracy Median Voting Threshold Experiments

From this, we see that a threshold value of 0.5 for the median voting performs marginally better, but not significantly

enough to make a real difference.

5. Analysis

From our experiments we concluded that our second approach to feature generation and classification was far more successful than our first approach. We believe that the main issue with our first approach was insufficient data. We were only able to generate about 300 total clips, which meant we had about 210 feature vectors to train on and 90 to test on. This simply was not enough data to train a robust classifier.

When we trained the model on all the frames, we had a much larger dataset to train on, as each clip would give us 60 feature vectors. This gives the model more of a chance to learn patterns in the data to make its predictions.

A deeper dive into our second approach to feature vector generation shows us that a threshold of 0.50 for the majority vote performs the best, and a median vote goes not significantly change the performance.

6. Conclusion

References

- [1] Zhe Cao, Gines Hidalgo, Tomas Simon, Shih-En Wei, and Yaser Sheikh. Openpose: Realtime multi-person 2d pose estimation using part affinity fields, 2018. 1
- [2] Datagen. Pose estimation: Concepts, techniques, and how to get started, 2020. 3
- [3] Anna Veronika Dorogush, Andrey Gulin, Gleb Gusev, Nikita Kazeev, Liudmila Ostroumova Prokhorenkova, and Aleksandr Vorobev. Fighting biases with dynamic boosting. *CoRR*, abs/1706.09516, 2017. 4
- [4] Theodoros Evgeniou and Massimiliano Pontil. Support vector machines: Theory and applications. volume 2049, pages 249–257, 09 2001. 4
- [5] Haoshu Fang, Shuqin Xie, and Cewu Lu. RMPE: regional multi-person pose estimation. *CoRR*, abs/1612.00137, 2016. 2
- [6] Kaiming He, Georgia Gkioxari, Piotr Dollár, and Ross Girshick. Mask r-cnn, 2017. 2
- [7] Rong Ji. Research on basketball shooting action based on image feature extraction and machine learning. *IEEE Access*, 8:138743–138751, 2020. 1, 2
- [8] Tsung-Yi Lin, Piotr Dollár, Ross Girshick, Kaiming He, Bharath Hariharan, and Serge Belongie. Feature pyramid networks for object detection, 2016. 2
- [9] Fionn Murtagh. Multilayer perceptrons for classification and regression. *Neurocomputing*, 2(5):183–197, 1991. 4
- [10] Leonid Pishchulin, Eldar Insafutdinov, Siyu Tang, Bjoern Andres, Mykhaylo Andriluka, Peter Gehler, and Bernt Schiele. Deepcut: Joint subset partition and labeling for multi person pose estimation, 2015. 2
- [11] Joseph Redmon, Santosh Divvala, Ross Girshick, and Ali Farhadi. You only look once: Unified, real-time object detection. In *2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 779–788, 2016. 1, 2

- [12] Alexander Toshev and Christian Szegedy. DeepPose: Human pose estimation via deep neural networks. In *2014 IEEE Conference on Computer Vision and Pattern Recognition*. IEEE, jun 2014. [2](#)
- [13] Viso.ai. Object detection in 2022: The definitive guide, 2020. [2](#)
- [14] Viso.ai. Human pose estimation with deep learning the ultimate overview, 2022. [1](#)